

Logistic Regression on LOAD DIGITS dataset from scratch

importing libraries

```
In [1]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 from sklearn.model_selection import train_test_split
        4 from keras.utils import np_utils
        5 import sklearn
```

Using TensorFlow backend.

Loading Dataset

```
In [2]: 1 from sklearn.datasets import load_digits
        2 digits=load_digits()
```

```
In [3]: 1 Y=digits['target']
        2 images=digits['images']
```


One hot Encoding

```
In [4]: 1 Y = np_utils.to_categorical(Y,10)
```

```
In [5]: 1 images=images/16
```

```
In [6]: 1 plt.imshow(images[99])
        2 Y[99]
```

Out[6]: array([0., 1., 0., 0., 0., 0., 0., 0., 0., 0.], dtype=float32)



Flattening

```
In [7]: 1 images_new=images.reshape(images.shape[0],-1)
        2 print(images_new.shape,Y.shape)
```

(1797, 64) (1797, 10)

Train Test Split

```
In [8]: 1 train_X,test_X,train_Y,test_Y=train_test_split(images_new,Y,test_size=0.3)
```

```
In [9]: 1 train_X=train_X.T
        2 test_X=test_X.T
        3 train_Y=train_Y.T
        4 test_Y=test_Y.T
```

Randomly intializing weights and bais

```
In [15]: 1 W=np.random.randn(train_X.shape[0],train_Y.shape[0])
        2 b=np.zeros((train_Y.shape[0],1))
```

```
In [16]: 1 print(W.shape)
2 print(b.shape)

(64, 10)
(10, 1)
```

LOGISTIC REGRESSION SGD with momentum

```
In [17]: 1 v1 = 0
2 v2 = 0
3 mu = 0.99
4 learning_rate = 1
```

```
In [20]: 1 for i in range(2000):
2     m=train_X.shape[1]
3     train_A=1/(1+np.exp(-(np.dot(W.T,train_X)+b)))
4
5     # column normalization
6     for j in range(train_X.shape[1]):
7         train_A[:,j]=train_A[:,j]/np.sum(train_A[:,j])
8     cost1=1/m*(- np.sum(train_Y* np.log(train_A+1e-6)))
9     print('iteration=',i,'-----train logloss-----',cost1)
10    #cost=sklearn.metrics.Log_loss(train_Y,A)
11
12    test_A=1/(1+np.exp(-(np.dot(W.T,test_X)+b)))
13    for j in range(test_X.shape[1]):
14        test_A[:,j]=test_A[:,j]/np.sum(test_A[:,j])
15    cost2=1/m*(- np.sum(test_Y* np.log(test_A+1e-6)))
16    print('iteration=',i,'-----test logloss-----',cost2)
17
18    dW=(1/m)*np.dot(train_X,(train_A-train_Y).T)
19    db=(1/m)*np.sum(train_A-train_Y)
20
21    v1 = mu * v1 - learning_rate * dW # integrate velocity
22    W += v1 # integrate position
23
24    v2 = mu * v2 - learning_rate * db # integrate velocity
25    b += v2 # integrate position
iteration= 1887 -----train logloss----- 1.498059368270154
iteration= 1887 -----test logloss----- 0.624485425594043
iteration= 1888 -----train logloss----- 1.498059368270154
iteration= 1888 -----test logloss----- 0.6244854156208702
iteration= 1889 -----train logloss----- 1.4980593665411939
iteration= 1889 -----test logloss----- 0.6244854056606044
iteration= 1890 -----train logloss----- 1.4980593648084026
iteration= 1890 -----test logloss----- 0.6244853957132129
iteration= 1891 -----train logloss----- 1.4980593630715806
iteration= 1891 -----test logloss----- 0.6244853857786635
iteration= 1892 -----train logloss----- 1.498059361330566
iteration= 1892 -----test logloss----- 0.6244853758569244
iteration= 1893 -----train logloss----- 1.4980593595852219
iteration= 1893 -----test logloss----- 0.6244853659479646
iteration= 1894 -----train logloss----- 1.4980593578354406
iteration= 1894 -----test logloss----- 0.6244853560517526
iteration= 1895 -----train logloss----- 1.4980593560811661
iteration= 1895 -----test logloss----- 0.6244853461682581
iteration= 1896 -----train logloss----- 1.4980593543223497
iteration= 1896 -----test logloss----- 0.6244853362974512
iteration= 1897 -----train logloss----- 1.498059352580014
```

```
In [21]: 1 #checking accuracy
2 count=0
3 for i in range(test_Y.shape[1]):
4     if(np.argmax(test_Y[:,i])==np.argmax(test_A[:,i])):
5         count+=1
6 accuracy=count/test_Y.shape[1]*100
7 print(accuracy)

25.555555555555554
```

LOGISTIC REGRESSION WITH ADAM

```
In [22]: 1 eps = 1e-8
2 beta1 = 0.9
3 beta2 = 0.999
4
5 learning_rate = 0.001
6 m1 = 0
7 v1 = 0
8
9 m2 = 0
10 v2 = 0
```

```
In [23]: 1 W=np.random.randn(train_X.shape[0],train_Y.shape[0])
2 b=np.zeros((train_Y.shape[0],1))
```

```
In [25]: 1 for i in range(2000):
2     m=train_X.shape[1]
3     train_A=1/(1+np.exp(-(np.dot(W.T,train_X)+b)))
4
5     # column normalization
6     for j in range(train_X.shape[1]):
7         train_A[:,j]=train_A[:,j]/np.sum(train_A[:,j])
8     cost1=1/m*(- np.sum(train_Y* np.log(train_A+1e-6)))
9     print('iteration=',i,'-----train logloss-----',cost1)
10    #cost=sklearn.metrics.Log_loss(train_Y,A)
11
12    test_A=1/(1+np.exp(-(np.dot(W.T,test_X)+b)))
13    for j in range(test_X.shape[1]):
14        test_A[:,j]=test_A[:,j]/np.sum(test_A[:,j])
15    cost2=1/m*(- np.sum(test_Y* np.log(test_A+1e-6)))
16    print('iteration=',i,'-----test logloss-----',cost2)
17
18    dW=(1/m)*np.dot(train_X,(train_A-train_Y).T)
19    db=(1/m)*np.sum(train_A-train_Y)
20
21    m1 = beta1*m1 + (1-beta1)*dW
22    mt1 = m1 / (1-beta1**(i+1))
23    v1= beta2*v1 + (1-beta2)*(dW**2)
24    vt1 = v1 / (1-beta2**(i+1))
25    W += - learning_rate * mt1 / (np.sqrt(vt1) + eps)
26
27    m2 = beta1*m2 + (1-beta1)*db
28    mt2 = m2 / (1-beta1**(i+1))
29    v2= beta2*v2 + (1-beta2)*(db**2)
30    vt2 = v2 / (1-beta2**(i+1))
31    b += - learning_rate * mt2 / (np.sqrt(vt2) + eps)
```

```
iteration= 1986 -----test logloss----- 0.5718389543856418
iteration= 1987 -----train logloss----- 1.3628685843306143
iteration= 1987 -----test logloss----- 0.5718513251534759
iteration= 1988 -----train logloss----- 1.3628896401197883
iteration= 1988 -----test logloss----- 0.5718637379920413
iteration= 1989 -----train logloss----- 1.3629107961635811
iteration= 1989 -----test logloss----- 0.5718761928396787
iteration= 1990 -----train logloss----- 1.3629320523204518
iteration= 1990 -----test logloss----- 0.5718886896347841
iteration= 1991 -----train logloss----- 1.3629534084490234
iteration= 1991 -----test logloss----- 0.5719012283158214
iteration= 1992 -----train logloss----- 1.3629748644081605
iteration= 1992 -----test logloss----- 0.5719138088213548
iteration= 1993 -----train logloss----- 1.3629964200568494
iteration= 1993 -----test logloss----- 0.5719264310899979
iteration= 1994 -----train logloss----- 1.3630180752542294
iteration= 1994 -----test logloss----- 0.5719390950604262
iteration= 1995 -----train logloss----- 1.363039829859652
iteration= 1995 -----test logloss----- 0.571951800671404
iteration= 1996 -----train logloss----- 1.363061683732624
```

```
In [26]: 1 #checking accuracy
2 count=0
3 for i in range(test_Y.shape[1]):
4     if(np.argmax(test_Y[:,i])==np.argmax(test_A[:,i])):
5         count+=1
6 accuracy=count/test_Y.shape[1]*100
7 print(accuracy)
```

88.70370370370371

REFERENCES :

- <http://cs231n.github.io/neural-networks-3/> (<http://cs231n.github.io/neural-networks-3/>)
- <https://stats.stackexchange.com/questions/219241/gradient-for-logistic-loss-function> (<https://stats.stackexchange.com/questions/219241/gradient-for-logistic-loss-function>)

OBSERVATION :

- Accuracy of SGD with momentum on test dataset is 25.5 %.
- Accuracy of ADAM on test dataset is 88.7 %.